

HARI KRISHNA BEKKAM

 portfolio |  HariKrishna070 |  HARI KRISHNA BEKKAM
 harikrishna.bekkam01@gmail.com |  +91 8639669877

SUMMARY

AI Engineer specializing in Agentic workflows, LLM optimization, and autonomous multi-agent systems. Proven track record in improving inference throughput and building production-ready GenAI applications.

WORK EXPERIENCE

Specialist Programmer - Infosys

Oct 2025 - Present

- Developing **Agentic AI solutions** to analyze legacy **VB.NET applications**, understand business logic and dependencies, and automate migration of legacy components into **.NET Core** and **Java Spring Boot services**.
- Developed automation scripts to analyze **ASP.NET (.aspx) files**, map legacy UI components to corresponding **Angular components**, and extract **business and database logic**, generating structured documentation and contextual knowledge to guide AI agents in accurate application migration.
- Designed **AI-assisted migration workflows** that leverage extracted application knowledge, business rules, UI mappings, and database dependencies to enable more accurate and context-aware legacy code transformation.

Agentic AI Engineer - AgentAnalytics.AI

July 2024 – Sep 2025

- Architected and productionized **multi-agent RAG applications** for enterprise clients using **LangChain**, **LangGraph**, **CrewAI**, and **AutoGen**, with **Langfuse** for end-to-end tracing, evaluation, and observability of agent workflows.
- Designed and developed **WaveFlow Studio**, a proprietary enterprise **Multi-Agent SDK** enabling scalable agentic applications through agent orchestration, memory management, and **MCP-based tool integration**.
- Partnered directly with **enterprise clients** to understand business requirements, architect tailored multi-agent AI solutions using WaveFlow Studio, and drive solutions from technical design through successful client adoption.
- Engineered and deployed **production-grade Multi-Agent RAG systems on AWS**, leveraging **SageMaker**, **EC2**, **ECR**, and **Docker** for scalable model serving, containerized deployments, and reliable AI workloads.
- Led **LLM inference optimization for CPU environments**, benchmarking **llama.cpp** and **vLLM** and tuning model configurations and runtime parameters to achieve **up to 300 tokens/sec** while improving deployment efficiency.

PROJECTS

ConvoTree – Agentic AI Conversation & Cognitive Memory Platform

- Engineered an **agentic conversation system** with deterministic context resolution, **Mem0 hierarchical memory**, and branch-aware reasoning for isolated AI workflows.
- Built an **async Python/FastAPI backend** using **SQLAlchemy**, recursive **SQL CTEs**, **tiktoken**, hierarchical summarization, and **SSE-based LLM streaming**.
- Developed a **multi-provider LLM orchestration layer** integrating **OpenAI**, **Anthropic**, **Gemini**, and **Ollama** through a standardized asynchronous inference interface.
- Built a **React + TypeScript AI interface** with **Zustand** and an interactive **SVG conversation-tree UI** supporting branching, context navigation, real-time streaming, and Markdown/code rendering.

EDUCATION

2021 - 2025 BTech(AI&DS) at **Vishnu Institute of Technology**

(CGPA: 8.9/10)

2019 - 2021 Intermediate at **Narayana Junior College**

(MARKS : 950/1000)

CERTIFICATIONS

IBM certified in the course of **Python for Data Science, AI & Development**.

IBM certified in the course of **Relational Databases (RDBMS)**.

NPTEL certified in the course of **Cloud Computing**.

SKILLS

Programming Languages :	Python, Data Structures and Algorithms
Data Analysis & Statistics :	Numpy, Pandas, Statistics
Machine Learning :	Supervised Learning, Unsupervised Learning, Feature Engineering
GenAI :	VectorDB, LLMs, RAG, AI Agents, Agentic AI, LangGraph, LangChain, Langfuse
Frameworks & Tools :	FastAPI, Docker, Git, GitHub
Databases :	SQL, MongoDB